

# Deep Learning - lab 6

Keras callbacks and CNNs

---

Prof. Stefano Carrazza

University of Milan and INFN Milan

# Keras Callbacks

---

A callback is an object that can perform actions at various stages of training (e.g. at the start or end of an epoch, before or after a single batch, etc).

You can use callbacks to:

- Periodically save your model to disk
- Do early stopping
- Write TensorBoard logs after every batch of training to monitor your metrics
- Get a view on internal states and statistics of a model during training
- ...and more

## Example: storing model parameters during training

```
# save the weights manually  
model = create_model()  
model.fit(...)  
model.save_weights('./checkpoints/my_checkpoint')  
  
# restore weights manually  
model = create_model()  
model.load_weights('./checkpoints/my_checkpoint')
```

# Keras callback API

## Example: storing model parameters during training

```
from tensorflow import keras
```

```
model = ... # define your model as usual
```

```
# declare callbacks
```

```
my_callbacks = [  
    keras.callbacks.ModelCheckpoint(filepath='model.{epoch}-{val_loss:.2f}.h5'),  
]
```

```
model.fit(data, epochs=10, callbacks=my_callbacks)
```

## ModelCheckpoint

**ModelCheckpoint** class


```
tf.keras.callbacks.ModelCheckpoint(  
    filepath,  
    monitor="val_loss",  
    verbose=0,  
    save_best_only=False,  
    save_weights_only=False,  
    mode="auto",  
    save_freq="epoch",  
    options=None,  
    initial_value_threshold=None,  
    **kwargs  
)
```

Callback to save the Keras model or model weights at some frequency.

**ModelCheckpoint** callback is used in conjunction with training using `model.fit()` to save a model or weights (in a checkpoint file) at some interval, so the model or weights can be loaded later to continue the training from the state saved.

# Keras callback API

## Example: early stopping



```
from tensorflow import keras

model = ... # define your model

# declare callbacks
my_callbacks = [
    keras.callbacks.EarlyStopping(patience=2),
]

model.fit(data, epochs=10,
          callbacks=my_callbacks)
```

## EarlyStopping

EarlyStopping class

```
tf.keras.callbacks.EarlyStopping(  
    monitor="val_loss",  
    min_delta=0,  
    patience=0,  
    verbose=0,  
    mode="auto",  
    baseline=None,  
    restore_best_weights=False,  
)
```

Stop training when a monitored metric has stopped improving.

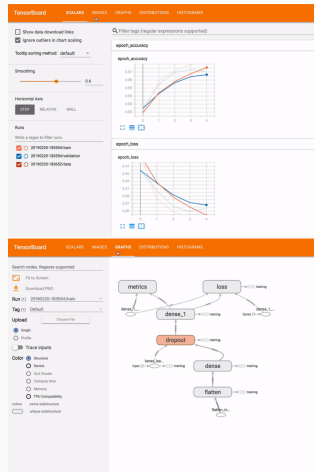
Assuming the goal of a training is to minimize the loss. With this, the metric to be monitored would be 'loss', and mode would be 'min'. A `model.fit()` training loop will check at end of every epoch whether the loss is no longer decreasing, considering the `min_delta` and `patience` if applicable. Once it's found no longer decreasing, `model.stop_training` is marked True and the training terminates.



# Keras callback API

**TensorBoard** provides the visualization and tooling needed for machine learning experimentation:

- Tracking and visualizing metrics such as loss and accuracy
- Visualizing the model graph
- Viewing histograms of weights, biases, or other tensors as they change over time
- Projecting embeddings to a lower dimensional space
- Displaying images, text, and audio data
- Profiling TensorFlow programs



# Keras callback API

## Example: storing logs for TensorBoard

```
from tensorflow import keras

model = ... # define your model

# declare callbacks
my_callbacks = [
    keras.callbacks.TensorBoard(log_dir='./logs'),
    keras.callbacks.EarlyStopping(patience=2),
]

model.fit(data, epochs=10, callbacks=my_callbacks)

# open tensorboard with
# tensorboard --logdir=path_to_your_logs
```

## TensorBoard

TensorBoard class

```
tf.keras.callbacks.TensorBoard(  
    log_dir="logs",  
    histogram_freq=0,  
    write_graph=True,  
    write_images=False,  
    write_steps_per_second=False,  
    update_freq="epoch",  
    profile_batch=0,  
    embeddings_freq=0,  
    embeddings_metadata=None,  
    **kwargs  
)
```

Enable visualizations for TensorBoard.

TensorBoard is a visualization tool provided with TensorFlow.


This callback logs events for TensorBoard, including:

- Metrics summary plots
- Training graph visualization
- Weight histograms
- Sampled profiling

# Convolutional Neural Nets

---

# Convolutional base



```
model = models.Sequential()  
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), activation='relu'))  
model.add(...)
```

## Conv2D layer

Conv2D class

```
tf.keras.layers.Conv2D(  
    filters,  
    kernel_size,  
    strides=(1, 1),  
    padding="valid",  
    data_format=None,  
    dilation_rate=(1, 1),  
    groups=1,  
    activation=None,  
    use_bias=True,  
    kernel_initializer="glorot_uniform",  
    bias_initializer="zeros",  
    kernel_regularizer=None,  
    bias_regularizer=None,  
    activity_regularizer=None,  
    kernel_constraint=None,  
    bias_constraint=None,  
    **kwargs  
)
```

2D convolution layer (e.g. spatial convolution over images).

This layer creates a convolution kernel that is convolved with the layer input to produce a tensor of outputs. If `use_bias` is True, a bias vector is created and added to the outputs. Finally, if `activation` is not `None`, it is applied to the outputs as well.

## MaxPooling2D layer

MaxPooling2D class

```
tf.keras.layers.MaxPooling2D(  
    pool_size=(2, 2), strides=None, padding="valid", data_format=None, **kwargs  
)
```

Max pooling operation for 2D spatial data.

Downsamples the input along its spatial dimensions (height and width) by taking the maximum value over an input window (of size defined by `pool_size`) for each channel of the input. The window is shifted by `strides` along each dimension.

The resulting output, when using the "valid" padding option, has a spatial shape (number of rows or columns) of:  $\text{output\_shape} = \text{math.floor}((\text{input\_shape} - \text{pool\_size}) / \text{strides}) + 1$  (when  $\text{input\_shape} \geq \text{pool\_size}$ )

The resulting output shape when using the "same" padding option is:  $\text{output\_shape} = \text{math.floor}((\text{input\_shape} - 1) / \text{strides}) + 1$

## Example: CNN classifier

```
num_classes = len(class_names)

model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])
```